

No Magic Required

How Zero-Knowledge Trust actually works — and why every piece has been battle-tested for years

February 2026

The question you should be asking

If someone tells you they have built a system where your secrets are completely protected, nobody can access your data, the architecture eliminates entire categories of attack, and the user experience is simpler than what exists today — you should be skeptical. The history of security is littered with systems that sounded too good to be true because they were.

So here is the honest answer to the honest question: how does Zero-Knowledge Trust actually work? What makes it real and not snake oil?

There is no trick. There is no novel cryptographic algorithm. There is no breakthrough mathematical discovery. There is no magic.

There is an architecture. ZKT is a pattern — a set of principles for how cryptographic primitives should be assembled so that the user holds the only keys to their digital life. VettID is an implementation of that pattern, and the one we will use as a concrete example throughout this document. But the pattern itself is open: any implementation that follows the same principles, using the same battle-tested primitives, would produce the same security properties.

The iPhone did not invent the touchscreen, the camera, or GPS. It assembled them in a way that created something new. ZKT does the same thing with cryptographic primitives. The innovation is the assembly, not any individual piece.

Every piece of Zero-Knowledge Trust has been beaten on by the best cryptographers in the world for years. The innovation is not any individual piece. It is the assembly.

The building blocks

Before we explain how the system works, here is everything it is built from. This is the inventory for VettID's implementation specifically, but any ZKT implementation would draw from the same pool of established primitives. Every item on this list is a published, peer-reviewed, widely deployed standard. None of them were invented by VettID.

Primitive	What it does	Who else uses it
X25519	Key exchange — two parties establish a shared secret without transmitting it	Signal, WhatsApp, WireGuard, TLS 1.3, SSH, Cloudflare
XChaCha20-Poly1305	Authenticated encryption — encrypts data and proves it has not been tampered with	Google, Cloudflare, 1Password, NordVPN, WireGuard
Ed25519	Digital signatures — proves a message came from a specific party	SSH, Signal, Tor, Solana, FIDO2/WebAuthn
Argon2id	Password hashing — makes brute force computationally infeasible (60 million times slower than SHA256 on GPUs)	Winner of the 2015 Password Hashing Competition; recommended by OWASP
HKDF	Key derivation — derives multiple purpose-specific keys from a single secret	TLS 1.3, Signal Protocol, WireGuard
TLS 1.3	Transport encryption — secures data in transit between devices	Every HTTPS website, every major browser, every cloud provider
AWS Nitro Enclaves	Hardware isolation — creates a sealed compute environment even AWS cannot access	AWS production infrastructure, Signal server-side processing
HMAC-SHA256	Integrity verification — detects any modification to stored data	JWT, TLS, IPsec, every major API authentication system

That is the complete list. There is no secret ingredient. Every algorithm, every protocol, every hardware component is published, documented, and available for independent review.

Built to evolve

Because ZKT is an assembly of standard primitives rather than a proprietary algorithm, it is designed to evolve as cryptography evolves. When post-quantum cryptography (PQC) algorithms reach production readiness — lattice-based key exchange replacing X25519, for example — a ZKT implementation can swap in the new primitives without changing the architecture. The pattern stays the same: encrypt, authenticate, use, rotate, re-encrypt. Only the specific algorithms change.

This is a direct benefit of building on standards rather than invention. A proprietary algorithm would require replacing the entire system when quantum computing advances. A standards-based assembly replaces individual components while the architecture continues to work. The building blocks are modular by design.

What the user actually does

Before we explain the cryptography, here is what ZKT feels like to a real user. We will use VettID's implementation as the example, but the mental model applies to any ZKT system. Because if the security only works in theory and the user experience is terrible, none of it matters. The mental model has four steps.

1. Enroll

You install the VettID app and scan a QR code from your account portal. The app walks you through creating a PIN and a password. Behind the scenes, your vault is being provisioned inside a hardware-isolated enclave, cryptographic keys are being generated, and your Protean Credential is being created. You see none of this. You see: scan code, create PIN, create password, done.

2. Add data

You add the things you want to protect: private keys, credentials, seed phrases, personal information. These go into your vault, encrypted with keys that only the enclave holds. Your phone stores an encrypted blob it cannot open. The vault stores the rest in an encrypted database. You manage what goes in through the app.

3. Connect

You connect to people and services. Each connection is a tendril — a scoped, revocable link between your vault and something outside it. You control the scope, the duration, and the access level. Connections are visible on your vine. You can see every active connection and sever any of them at any time.

4. Go

That is it. Your vault handles the cryptography. Your vine manages the connections. Your secrets are used when needed and never exposed. You open the app, enter your PIN, and your digital life is at your fingertips. When you need to authorize something sensitive, you enter your password. The system works.

Enroll. Add data. Connect. Go. That is the entire user experience. Everything else is handled by the architecture.

Organize your vine, don't overload your root

One critical design principle makes the system work well in practice: keep your root lean. Your vault is a keychain, not a filing cabinet. It holds the

keys to your digital life, not the data itself. And your connections should be organized thoughtfully, not piled on top of one another.

What not to do: Don't add every service, every contact, and every credential as a direct connection from the root. If your root has two hundred connections dangling off it, you have recreated the problem ZKT was designed to solve — a flat, unmanageable tangle where everything has the same level of access and visibility.

What to do instead: Think of your vine the way you think about files on a computer. You do not put every file on the desktop. You organize into folders. Your vine works the same way. Your root connects to a few high-trust branches: your financial identity, your professional identity, your personal connections. Each branch manages its own connections. Your bank tendril connects to financial services. Your professional tendril connects to work tools. Your personal tendril connects to family and friends.

This organization is not just aesthetics. It is security architecture. A compromise at one branch does not propagate to another. A misbehaving service connection on your professional branch cannot access your financial branch. The structure of your vine is your security posture, and a well-organized vine is a well-protected vine.

Your vault stores the cryptographic keys that control each branch. The data — your medical records, your financial accounts, your employment history — stays in the systems designed to manage it. Your vine carries authority, not data. The root stays lean regardless of how many connections you have, because the root holds keys, not content.

How it actually works

Now the machinery. We will walk through VettID's implementation step by step, naming every algorithm at each stage. Other ZKT implementations might choose different primitives from the same families — a different

authenticated encryption cipher, a different key exchange protocol — but the pattern would be the same. No hand-waving.

The Protean Credential

Your Protean Credential is an encrypted package that lives on your phone. It contains your identity keypair, your private keys, your seed phrases, your password hash, and your policies. You hold this package, but you cannot open it. It is encrypted with a Credential Encryption Key (CEK) that only your vault holds.

Encryption: The CEK is an X25519 keypair. The credential blob is encrypted using XChaCha20-Poly1305 authenticated encryption. The CEK private key — the only thing that can decrypt the blob — lives inside the vault's encrypted SQLite database, inside the Nitro Enclave. Both halves of the CEK keypair are held by the vault. Your phone never has the CEK private key. Your phone has the locked box. The vault has the key.

Two factors, two purposes

ZKT uses two separate authentication factors, each protecting a different layer. They are not interchangeable.

Vault PIN (6 digits): Entered when you open the app. The PIN is sent to the enclave supervisor, which combines it with sealed material from AWS KMS to derive the Data Encryption Key (DEK). The DEK decrypts the vault's SQLite database. Wrong PIN means wrong DEK, which means the database decryption fails cryptographically. The PIN is not compared against a stored value. It is verified by whether the math works. There is no hash to steal.

Credential password: Entered for each sensitive operation (signing a transaction, using a key). The password is hashed with Argon2id on your device, encrypted with a single-use User Transaction Key (UTK), and sent to the vault. The vault decrypts it with the corresponding Ledger Transaction Key (LTK), compares the hash to the one stored inside your Protean Credential, and proceeds only if they match. The UTK/LTK pair is single-use — destroyed after one operation — so even if an attacker intercepts the encrypted hash, they cannot replay it.

	Vault PIN	Credential Password	Why both?
Protects	Vault storage (DEK derivation)	Operation authorization	Compromising one does not compromise the other
When used	Each session (app open)	Each sensitive operation	PIN is convenience + security; password is per-action gate
Verified by	Crypto failure (wrong DEK = decryption fails)	Hash comparison (Argon2id)	PIN verified by hardware; password verified by software

What happens when you use a secret

You want to sign a Bitcoin transaction. Here is exactly what happens, step by step, algorithm by algorithm.

Step 1. Your app sends the encrypted Protean Credential blob and the operation request to the vault-manager inside the Nitro Enclave.

Step 2. The vault-manager decrypts the credential using the CEK private key (X25519 + XChaCha20-Poly1305). Your secrets now exist in hardware-isolated enclave memory. Nowhere else.

Step 3. The vault challenges you: enter your credential password. It tells the app which UTK to use for encrypting the response.

Step 4. You enter your password. The app hashes it with Argon2id (64 MB memory cost, 3 iterations, parallelism 4) and encrypts the hash with the specified UTK public key (X25519). The encrypted hash is sent to the vault.

Step 5. The vault decrypts the hash with the corresponding LTK private key. The LTK is single-use: it is destroyed immediately after decryption. The vault compares the received hash against the hash stored in the credential. If they match, you are authorized.

Step 6. The vault loads your Bitcoin private key from the credential, signs the transaction (secp256k1), and immediately zeros the key from memory. The key existed in enclave memory for milliseconds.

Step 7 — the protean part. The vault generates a brand new CEK keypair (X25519). It re-encrypts your entire credential with the new CEK public key. It stores the new CEK private key in the DEK-encrypted SQLite database and syncs to S3. It returns the re-encrypted credential blob, the transaction signature, and fresh UTKs to your app.

Step 8. Your app stores the new encrypted blob and the new UTKs. The old CEK is dead. The old blob is cryptographic garbage. If anyone intercepted it, they have nothing usable.

The credential changes shape after every single use. That is why it is called protean. There is no stable key to target. An attacker who steals today's blob has a package locked with a key that was destroyed the moment the operation completed.

X25519 key generation takes approximately 0.05 milliseconds. That is why rotation after every operation is practical. There is no performance penalty for changing every lock in the building after every use, because the locks are nearly free to manufacture.

The hardware guarantee

The obvious question: what if VettID is the attacker? What if the company that runs the vault modifies the code to exfiltrate secrets? What prevents a rogue employee, a compromised server, or a government subpoena from accessing your data?

The answer is hardware, not trust. A ZKT implementation requires a hardware-isolated environment where secrets can be processed without exposure. VettID uses AWS Nitro Enclaves. Other implementations might use different trusted execution environments — the principle is the same. The secrets must be processed inside a sealed compute environment that even the operator cannot access.

A Nitro Enclave is a hardware-isolated virtual machine with three critical properties: no persistent storage, no external network access, and no interactive access — not even by AWS administrators. The enclave

communicates with the outside world through a single constrained channel (vsock), and everything passing through that channel is encrypted.

Attestation: proving the code is honest

When your app connects to the vault, it does not take VettID's word that the correct code is running. It verifies cryptographically.

How attestation works: The enclave's Nitro Security Module (hardware, not software) generates an attestation document containing PCR values — cryptographic hashes of the exact code running inside the enclave. Your app sends a random nonce, which is included in the attestation document to prevent replay. The app verifies the AWS signature on the attestation, checks the PCR values against VettID's published values, and confirms the nonce matches. If any check fails, the app refuses to proceed.

This is not a policy. It is physics. The Nitro hardware generates the attestation. AWS cannot forge it. VettID cannot forge it. If the code has been modified by anyone for any reason, the PCR values change, and every app in the world will reject the connection.

Sealed storage: the code protects the data

Your vault's encryption material is sealed to the enclave's code identity using AWS KMS. The sealed material can only be unsealed by a Nitro Enclave running the exact same code (matching PCR values). If VettID deploys different code — even a one-byte change — the sealed material cannot be unsealed. The vault data becomes inaccessible.

What this means: VettID cannot deploy a modified enclave that exfiltrates data, because the modified enclave cannot unseal the existing vault data. The PCR values would differ. The KMS policy would reject the unseal request. The vault would refuse to start. Even the parent EC2 instance that hosts the enclave sees only encrypted blobs it cannot read. The math prevents exfiltration, not a policy document.

The parent process that manages network traffic and S3 storage explicitly cannot access vault keys, plaintext data, or session keys. It sees encrypted blobs and routing metadata. That is all.

You approve enclave updates

Enclave code changes over time — bug fixes, security patches, new features. When this happens, the new code has different PCR values. Your vault's sealed material is bound to the old PCR values and cannot be unsealed by the new code. This is a feature, not a bug. It means nobody can silently swap the code your vault runs.

In VettID's implementation, the process is transparent: when a new enclave version is deployed, your app is notified. The old enclave and the new enclave run simultaneously. You are not moved to the new version until you approve. When you do, your sealed material is re-sealed for the new PCR values inside the old enclave (which can still unseal it), and then verified on the new enclave before cutover. The old version stays available for rollback.

Why this matters: In the current model, service providers push server-side changes constantly and you have no visibility or control. You trust that every update is benign. In ZKT, a code change produces different PCR values, your app can detect it, and you decide when to move. The attestation chain makes invisible changes impossible and gives users meaningful control over the software that handles their secrets.

Compared to what?

ZKT sounds ambitious until you compare it to what currently passes for security. In the current model, your passwords exist in databases operated by hundreds of companies, each with hundreds of employees who have database access, each running infrastructure with thousands of potential

vulnerabilities. A single breach at any one of these companies exposes your credentials. This is not a theoretical risk. It is a weekly news story.

ZKT's claim is actually more modest than the current model's implicit claim.

Current model claims	ZKT claims
Hundreds of employees with database access will never make a mistake or act maliciously	The hardware enclave is isolated and the math is correct
Every server will remain uncompromised indefinitely	Even a compromised server sees only encrypted blobs it cannot read
Security policies will be followed perfectly by every person in every organization	Architectural properties hold regardless of who follows what policy
Bearer tokens will not be intercepted or stolen	Credentials change after every use; stolen credentials decrypt with dead keys

The current model requires trusting thousands of people and systems to behave correctly, indefinitely. ZKT requires trusting published algorithms and hardware isolation. One of these is a human problem with no known solution. The other is a mathematics problem with decades of validation.

ZKT does not claim to be perfectly secure. It claims to reduce the trust surface from “every person and system that touches your data” to “the math and the hardware.” That is not a bigger claim. It is a dramatically smaller one.

What could go wrong

No security system is perfect. Honest engineering requires acknowledging attack surfaces. Here is what we know about ZKT's.

Device compromise: If malware captures your PIN and password as you type them, and exfiltrates your encrypted credential blob before you use it, an attacker could race you to the vault. This is mitigated by device attestation (the vault verifies your device has not been rooted or tampered with) and by the single-credential design:

the first person to use the credential rotates the CEK, making the other copy worthless. But a sophisticated, targeted device attack remains a risk.

Enclave vulnerability: If a vulnerability is discovered in the Nitro Enclave hardware itself — not the code, but the isolation boundary — the security model would be affected. This is a risk shared by every system that relies on hardware security: Intel SGX, ARM TrustZone, Apple's Secure Enclave. Nitro has a strong track record, but no hardware is perfect.

Lost device: A ZKT implementation must address device loss. VettID handles this with a hosted backup: every time the vault issues a new credential blob to your phone, it simultaneously stores a copy in VettID's backend. This backup is the same encrypted blob — it is useless without a genuine Nitro Enclave and your PIN. If you lose your device, you log into the account portal, request recovery, and wait 24 hours (an intentional delay that gives you time to cancel if the request was unauthorized). After the waiting period, you scan a QR code on your new device and enter your PIN. The backup blob is delivered, the vault warms, and you are back. No seed phrase required. Other ZKT implementations might handle backup differently, but the principle is the same: the backup is encrypted material that cannot be used without the user's authentication factors and a trusted execution environment.

These are real risks. They are also dramatically smaller and more contained than the risks in the current model, where a single database breach at a single company can expose millions of users simultaneously. ZKT's failure modes are individual and targeted. The current model's failure modes are systemic and catastrophic. And for the most common risk — losing a device — the backup model means recovery is straightforward, while the encrypted backup itself remains useless to anyone who breaches the backend.

The theft detection race

One property of the single-credential design deserves its own explanation, because it is counterintuitive and often mistaken for a weakness. It is actually a deliberate security feature.

Because the credential changes after every use, there is only one valid copy at any given moment. If an attacker steals your encrypted credential blob, they have a narrow window: they must present it to the vault and authenticate before you do. The moment either party uses the credential, the CEK rotates and the other copy becomes garbage.

If you use your credential first (which is overwhelmingly likely, since you use it routinely), the attacker's stolen copy is worthless. If the attacker somehow uses it first, your next attempt fails, which is itself a detection signal — the vault knows something is wrong, because a credential was used that should not have been. This is not a bug. It is a tripwire.

In the current model, a stolen password works forever unless you change it, and you rarely know it has been stolen. In ZKT, a stolen credential works exactly once, and using it immediately alerts both the vault and the legitimate user.

The open challenge

If the primitives are standard and the architecture is sound, the appropriate response to skepticism is not persuasion. It is invitation.

Every cryptographic primitive used in ZKT is published. Every algorithm is named. In VettID's implementation, the enclave code produces verifiable PCR hashes that anyone can check. The architecture is documented in detail. Any ZKT implementation worthy of the name should be designed to be audited, not trusted on faith.

The invitation: Review the primitives. Audit the code. Verify the attestation chain. Try to find the flaw. If the architecture is broken, we want to know. If a primitive has a weakness we have not accounted for, we want to know. If the assembly introduces a vulnerability that the individual components do not have, we want to know. The best security is security that has been attacked by smart people and survived.

We do not ask you to trust us. We ask you to verify. That is the entire point of zero-knowledge trust: trust is replaced by verification. The same principle that protects your data from VettID protects your judgment from our marketing.

One more time, simply

Your vault is a keychain, not a filing cabinet. Your phone holds a locked box. The vault holds the key. When you need your secrets, the vault opens the box inside a hardware-sealed room, uses your secrets, locks everything back up with a new key, and hands you the new locked box. The old key is destroyed. Nobody else was in the room. Nobody can get into the room. The room's walls are guaranteed by hardware, not promises.

You enroll. You add your data. You connect to the world. You go. And you organize your vine thoughtfully — not every connection at the root, but branches that reflect how you actually live your digital life.

Every piece of this has been tested for years. Every algorithm is published. Every claim is verifiable.

Sufficiently well-assembled engineering is indistinguishable from magic. But it is not magic. It is X25519, XChaCha20-Poly1305, Argon2id, hardware enclaves, HKDF, and Ed25519 — arranged so that no one except you can access your digital life. VettID is one implementation. The pattern is open. The pieces are boring. The assembly is the breakthrough.